

Conceptual Modeling Language Specification

Version 1.0 (Draft)

Quenio Cesar Machado dos Santos
Universidade Federal de Santa Catarina*

November 2017

* Initially developed as part of the author's Bachelor Technical Report in Computer Sciences

Contents

I	Language and Compiler	1
1	The Language	2
2	The Compiler	3
2.1	The Frontend	4
2.2	The Backend	4
3	Specification and Notations	5
II	Conceptual Modeling	7
4	Concepts	8
4.1	Example	8
4.2	Syntax	8
4.3	Constraints	11
5	Properties	12
5.1	Example	12
5.2	Syntax	13
5.3	Constraints	14
6	Attributes	15
6.1	Example	15
6.2	Syntax	16
7	Derived Attributes	17
7.1	Example	17
7.2	Syntax	18

8	Associations	19
8.1	Example	20
8.2	Syntax	22
8.3	Constraints	22
8.4	Translation	23
9	Derived Associations	25
10	Generalization / Specialization	26
10.1	Example	26
10.2	Syntax	27
10.3	Constraints	27
11	Abstract Concepts	30
11.1	Example	30
11.2	Syntax	31
11.3	Constraints	31
III	Expressions	34
12	Expressions	35
12.1	Example	35
12.2	Syntax	35
13	Literal Values	39
14	Prefix Expressions	40
15	Infix Expressions	41
16	Conditional Expressions	42
17	Path Expressions	43
18	Query Expressions	44
IV	Type Checking	45
19	Types	46
20	Primitive Types	47

<i>CONTENTS</i>	iii
20.1 Example	47
20.2 Syntax	49
21 Sequence Types	51
V Code Generation	52
22 Templates	53
23 Constructors	54
24 Tasks	55
25 Targets	56
VI Organization and Sharing	57
26 Modules	58
27 Libraries	59
VII Appendices	60
A CML Concrete Syntax (Grammar)	61
A.1 ANTLR Grammar	62
B CML Abstract Syntax (Metamodel)	66
C CML Abstract Syntax Tree (Instantiation)	67
D CML Constraints (Validations)	68
E Language Specification Notation	70
Bibliography	71

Listings

4.1	Concept Examples	9
4.2	Concept Concrete Syntax	9
4.3	Concept AST Instantiation	10
4.4	Concept Constraints	11
5.1	Property Examples	12
5.2	Property Concrete Syntax	13
5.3	Property AST Instantiation	13
5.4	Property Constraints	14
6.1	Attribute Example	16
8.1	Association Example	21
8.2	Association Concrete Syntax	22
8.3	Association AST Instantiation	23
8.4	Association Constraints	24
11.1	Abstract Concept Example	32
11.2	Abstract Concept Constraints	33
12.1	Expression Example	36
12.2	Expression Concrete Syntax	37

Figures

2.1	An architectural overview of the CML compiler.	3
4.1	Concept Abstract Syntax	10
10.1	Generalization Examples	28
10.2	Generalization Constraints	29
12.1	Expression Abstract Syntax	38
20.1	Example of <i>Primitive Types</i>	48
20.2	Type Declaration Syntax	49
20.3	Type AST Instantiation	50

Tables

20.1 Core Primitive Types in CML.	49
20.2 Additional Primitive Types in CML.	50

Part I

Language and Compiler

One

The Language

This document specifies the *Conceptual Modeling Language*, or CML for short. CML enables the modeling of the information of software systems. It focuses on modeling the structural aspects of such systems, having less emphasis on the behavioral aspects. Using CML, it is possible to represent the information as understood by the system users, while disregarding its physical organization as implemented by target languages or technologies.

In this first part of the CML specification, the first chapter will provide an overview of the CML compiler's architecture, and the second chapter describes the organization and notation used in the remainder of this document. The second part describes that structural constructs of the language that enable conceptual modeling. The third part covers values and expressions. The fourth part focuses on the semantics of type checking. The fifth part describes code generation. The last part will cover organization and sharing of conceptual models.

Two

The Compiler

The CML compiler has as *input*, source files defined using its own conceptual language (as specified in this document), which provides an abstract syntax similar to (but less comprehensive than) a combination of UML [4] and OCL [3]; and, as *output*, any target languages based on extensible templates, which may be provided by the compiler's base libraries, by third-party libraries, or even by developers.

The CML compiler's overall architecture follows the standard compiler design literature [2]. An overview diagram of the architecture is shown in figure 2.1. The two main components of the compiler, and the




Figure 2.1: An architectural overview of the CML compiler.

artifacts they work with, are presented in the next sections.

2.1 The Compiler Frontend

The frontend receives as input the *CML source files*. It will parse the files and generate an internal representation of the *CML model*.

Syntactical and semantic validations will be performed at this point. Any syntax and constraint errors are presented to the developer, interrupting the progress to the next phase. If the *source files* are parsed and validated successfully, then the internal representation (the AST) of the *CML model* is provided as the input to the *backend* component.

2.2 The Compiler Backend

The backend receives the *CML model AST* as input. Based on the *target specification* provided by the AST, chooses which *extensible templates* to use for code generation. The *target files* are then generated, and become available to be consumed by other tools. The *task declaration* plays the key role of determining the kind of *target* to be generated.

CML extensible templates are implemented in StringTemplate [7]. The CML compiler uses StringTemplate for two purposes:

- *File names and directory structure*: each type of target generated by the CML compiler requires a different directory structure. The CML compiler expects each target type to define a template file named “files.stg” (also known as *files template*), which will contain the path of all files to be generated. The *files template* may use information provided by the *task declaration* (specified in chapter §25) in order to determine the file/directory names.
- *File content generation*: each file listed under the *files template* will have a corresponding *content template* that specifies how the file’s content must be generated. The *content template* will receive as input one root-level element of the CML model, which will provide information to generate the file’s content. Each type of top-level model element should have a corresponding *content template*. Templates are described in *Code Generation* part of this specification.

Three

Specification and Notations

The first two parts of this specification specify all elements of CML metamodel needed for conceptual modeling. Every chapter in these two parts starts with a definition, which is followed by: an example; the specification of the concrete syntax; the abstract syntax; and then how to transform the concrete syntax into the abstract one.

The definition of each CML metamodel element is stated in plain English on the first paragraph of the chapter. If a correspondence exists to an element of the Entity-Relationship (ER) [1] metamodel, or to an element of the Unified Modeling Language (UML) [4] metamodel, it is provided.

Examples are provided in a separate section for each declaration of a metamodel element. These sections refer to a `verbatim` figure containing the examples, and describes them as needed. The examples are provided for illustrative purposes only, and they are *not* intended to be normative. They may be excerpts of larger CML source files, and thus may not be successfully compiled on their own.

The concrete syntax of each CML metamodel element is described on its own section. This type of section refers to a `verbatim` figure, which contains the actual ANTLR [6] grammar specifying the syntax for the CML metamodel element in question, and it must be considered normative. The appendix § A presents all the grammar rules in a single listing.

The abstract syntax of each CML metamodel element is described on the next section. This type of section refers to two types of figure: the first figure presents a class diagram with the EMOF [5]-based metamodel of the element being described; the second figure specifies the transformation from the concrete syntax into instances of the metamodel classes, which are the nodes of the abstract syntax tree (the intermediate representation described in the *Compiler Overview*). The notation used to specify the transformations is presented in the appendix § E. Both figures

must be considered normative.

The constraints of each CML metamodel element are described on its own section. These sections refer to a verbatim figure, which contains the OCL [3] invariants (and its definitions) of the CML metamodel element in question, and it must be considered normative. Each invariant has a name in the format `inv_name` so that it can be referred by the compiler's error messages and users. Derived properties may also be defined before the constraints in order to simplify the constraint expressions. The appendix § D presents all the constraint rules in a single listing.

All metamodel elements referred by one of the descriptions defined above (definitions, examples, etc.) are emphasized in *italic*. If the descriptions of a CML metamodel element refer to another metamodel element, the corresponding chapter or section defining the other element is provided in parenthesis, like so (§ 3).

Some sections may not follow the structure defined above. These normally provide additional semantic information in plain English, which cannot be described using the notations presented above.

Part II

Conceptual Modeling

Four

Concepts

A *concept* in CML represents anything that has a coherent, cohesive and relevant meaning in a domain. In the ER [1] metamodel, it corresponds to an *entity set* (or an *entity type*); in UML [4], to a *class*. The CML *concept* differs, however, from the UML *class*, because it has only *properties* (§5), while the UML *class* may also have *operations*.

4.1 Example

Listing 4.1 presents some examples of *concepts* declared in CML. As shown, a *concept* may have zero or more *properties* (§5), and a *property* may optionally declare a *type* (§20, §21). Also, as shown in the concept **EBook** of the example, a *concept* may specialize (§10) another *concept*.

4.2 Syntax

Listing 4.2 specifies the syntax used to declare a *concept*. The **concept** keyword is followed by a NAME. Optionally, a list of other NAMES may be enumerated, referring to other *concepts* that are generalizations (§10) of the declared *concept*. A list of *properties* (§5) may be declared under the **concept** block. And the **abstract** keyword may precede the **concept** keyword, making a *concept* abstract (§11).

Figure 4.1 presents the *Concept* metaclass in an EMOF [5] class diagram, and listing 4.3 specifies the *concept* transformation from its concrete syntax to its abstract syntax. For each *concept* parsed by the compiler, an instance of the *Concept* class will be created, and its properties will be assigned according to parsed information:

```
// Empty concept:
@concept Book;

// Property without a type:
@concept TitledBook
{
    title;
}

// Property with the String type:
@concept StringTitledBook
{
    title: String;
}

// Specializing another concept:
@concept Ebook: Book;
```

Listing 4.1: Concept Examples

```
conceptDeclaration returns [Concept concept]:
    ABSTRACT? 'concept' NAME
    ( ':' generalizations )?
    ( ';' | propertyList );

generalizations:
    NAME ( ',' NAME ) *;
```

Listing 4.2: Concept Concrete Syntax

- *name*: assigned with the value of the terminal node NAME.
- *abstract*: set to *true* if the **abstract** keyword is found before the **concept** keyword; otherwise, set to *false*.
- *elements*: an *ordered set* referencing all *properties* parsed in the **concept** block.
- *generalizations*: an *ordered set* referencing all *concepts* whose NAMES were enumerated in the *GeneralizationList*.



Figure 4.1: Concept Abstract Syntax

```

node Concept:
  'abstract'?
  'concept' NAME
  ( ':' GeneralizationList )?
  ( ';' | PropertyList )
{
  name = NAME;
  abstract = 'abstract'?;
  elements = PropertyList.Property*;
  generalizations = for name in GeneralizationList.NAME*
    | yield Model.concept[name];
}

node GeneralizationList: NAME ( ',' NAME )*;

```

Listing 4.3: Concept AST Instantiation

```
context Concept

inv unique_concept_name:
  parent.concepts
    ->select(c | c != self and c.name = self.name)
    ->isEmpty()
```

Listing 4.4: Concept Constraints

4.3 Constraints

Listing 4.4 presents the invariants of the *Concept* metaclass:

- *unique_concept_name*: Each *concept* must have a unique NAME within its *module* (§26).

Five

Properties

A *property* in CML may hold *values* of *primitive types* (§20), in which case they correspond to *attributes* (§6) on the ER [1] and UML [4] metamodels. Or they may hold references, or *sequences* of references (§21), linking to instances of other *concepts* (§4), in which case they correspond to a *relationship* on the ER metamodel, or to *associations* (§8) on the UML metamodel.

5.1 Example

Listing 5.1 shows some *properties* declared in CML. As shown, a *property* may be an *attribute* (§6) of a *primitive type* (§20), or represent the *role* (or *end*) of an *association* (§8).

```
— Attributes of primitive types:
@concept Book
{
    title: String;
    quantity: Integer;
}

— Role in unidirectional association:
@concept Order
{
    customer: Customer;
}
```

Listing 5.1: Property Examples

```

propertyList:
    '{' (propertyDeclaration ';'*) '}'

propertyDeclaration returns [Property property]:
    DERIVED? NAME (':' typeDeclaration)? ('=' expression)?

DERIVED: '/'

```

Listing 5.2: Property Concrete Syntax

```

node PropertyList: '{' (Property ';')* '}'

node Property: '/'? NAME (':' Type)? ('=' STRING)?
{
    name = NAME;
    derived = '/'?;
    value = unwrap(STRING?);
    type = Type?;
}

```

Listing 5.3: Property AST Instantiation

5.2 Syntax

Listing 5.2 specifies the syntax used to declare a *property*. The NAME is followed by a *typeDeclaration* (§20 and §21). Optionally, an *expression* (§12) may be specified in order to set the initial value.

Figure ?? presents the *Property* metaclass in an EMOF [5] class diagram of the CML metamodel, and listing 5.3 specifies the transformation from the *property* concrete syntax to its abstract syntax. For each *property* parsed by the compiler, an instance of the *Property* class will be created, and its properties will be assigned according to parsed information:

- *name*: assigned with the value of the terminal node NAME.
- *type*: if *typeDeclaration* is provided, *type* is set with the instance of the *Type* class matching the *typeDeclaration*.
- *expression*: if provided, it contains the instance of the *Expression* class matching the parsed *expression*.

```

context Property
inv unique_property_name:
  self.scope.properties
    ->select(p | p != self and p.name = self.name)
    ->isEmpty()

context Property
inv property_type_specified_or_inferred:
  type->notEmpty() or expression->notEmpty()

context Property
inv property_type_assignable_from_expression_type:
  type->notEmpty() and expression->notEmpty() implies
    type.isAssignableFrom(expression.type)

```

Listing 5.4: Property Constraints

5.3 Constraints

Listing 5.4 presents the invariants of the *Property* metaclass:

- *unique_property_name*: Each *property* must have a unique NAME within its *concept* (§4).
- *property_type_specified_or_inferred*: Either the *property* explicitly defines a *type* or it defines an *expression*, from which the type is inferred. That is required for both regular, slot-based *properties* (which may provide an *initialization expression*) and *derived properties* (which may have an *expression* defining the derivation).
- *property_type_assignable_from_expression_type*: When both a *type* and *expression* are defined for a *property*, the *type* inferred from the *expression* should be assignable to the declared *type*. That is required for both regular, slot-based *properties* (which may provide an *initialization expression*) and *derived properties* (which may have an *expression* defining the derivation).

Six

Attributes

In CML, *attributes* are *properties* (§5) of *primitive types* (§20). They correspond to the *Attribute* metaclass in the ER [1] metamodel; in the UML [4] metamodel, to the association *attribute* between the metaclass *Class* and the metaclass *Property*.

Attributes serve as a *slot* that holds a value of the specified *primitive type*. An initial value may be specified as an *expression* (§12). Some *attributes*, however, may be continuously derive their *value* from an *expression* (not only initially), in which case they are called *derived attributes* (§7).

While initial values are only set when a *concept* (§4) is instantiated, the value of *derived attributes* is always evaluated from the given *expression*, and they cannot be set any other way.

6.1 Example

Listing 6.1 presents some examples of *attributes* declared in CML. As shown, the attribute **a** is a regular attribute definition that specifies the *primitive type* (§20) of the values that can be held by the *attribute's* slot. The attribute **b** is an example showing how an *attribute* can be defined with an initial value. As shown by the attribute **c**, an attribute may be derived from an *expression* that refers to other *attributes*. In order to differentiate *attributes* with initial values from *derived attributes*, a forward slash ("/") prefixes the name of the latter. Attributes **d** and **e** are examples where the type of the attribute, instead of being specified, is inferred from the given *expression*. Type inference is possible for both regular, slot-based *attributes* and *derived attributes* that provide an *expression*.

```

@concept Attributes
{
  — Attribute with a slot for values of a primitive type:
  a: Integer;

  — An attribute with an initial value:
  b: String = "initial_value";

  — A derived attribute:
  /c: Decimal = 2.0 * a;

  — An attribute with type inferred from its initial value:
  d = 3; — Inferred as Integer based on constant "3"

  — A derived attribute with type inferred from expression:
  /e = 2.0d * a; — Inferred as Double based on "2.0d * a"
}

```

Listing 6.1: Attribute Example

6.2 Syntax

Listing 5.2 specifies the syntax used to declare any kind of *property* (§5), including *attributes*. The NAME of an *attribute* is followed by a *typeDeclaration* of a *primitive type* (§20). Optionally, an *expression* (§12) may be specified in order to set the initial value. A *derived attribute* must be prefixed with the forward-slash character, as specified by DERIVED, in which case the given *expression* defines the value of the *attribute* at all times.

Since an *attribute* in CML is just a *property* (§5) with *primitive types* (§20), the *property* metaclass in the CML metamodel is used to represent *attributes*. Figure ?? presents the *property* metaclass in an EMOF [5] class diagram, and listing 5.3 specifies the *property* transformation from its concrete syntax to its abstract syntax.

Seven

Derived Attributes

A *concept* in CML may have *attributes* (§6) that do not hold specific *values*, but instead provide a *value* derived from an *expression* (§12). These are called *derived attributes*. Unlike an *expression* used to initialize a *non-derived attribute*, the expression of a derived attribute is evaluated every time the attribute is queried.

In the UML [4] metamodel, the *Property* metaclass has a meta-attribute named *isDerived*, which determines whether an *attribute* is derived or not. A *derived attribute* in UML may be defined using a OCL [3] constraint; while CML has *expressions* as part of the language.

The ER [1] metamodel, in its original form, does not allow for the differentiation of *derived attributes* as part of an *entity set*, but it is possible to define *retrieval operations* whose results would equal to *values* of *derived properties* in CML. It can be said, however, that ER, by defining an *attribute* as a function from the *entity set* to the *value set*, does not prescribe that all *attributes* are memory-based, nor does it prevent the definition of an *attribute function* as an *expression*.

The CML metamodel and its syntax, on the other hand, define whether an *attribute* is memory-based (a *non-derived attribute*) or it is derived from an *expression* (a *derived attribute*).

7.1 Example

Figure ?? presents two examples of *derived attributes* declared in CML. As shown, the attribute **c** is derived from an *expression* that refers to other *attributes*. In order to differentiate *attributes* with initial values, such as **b**, from *derived attributes*, such as **c**, a forward slash ("/") prefixes the name of the latter. The attribute **e** is an example of a *derived attribute* where the type is inferred from the given *expression*, instead of being specified.

7.2 Syntax

Figure ?? specifies the syntax used to declare any kind of *property* (§5), including *derived attributes*. A *derived attribute* must be prefixed with the forward-slash character, as specified by DERIVED, in which case the given *expression* provides the value of the *attribute* every time it is fetched.

The *property* metaclass in the CML metamodel is used to represent *attributes*. Figure ?? presents the *property* metaclass in an EMOF [5] class diagram, and figure ?? specifies the *property* transformation from its concrete syntax to its abstract syntax. The *derived* property of the *Property* metaclass defines whether the *attribute* is derived or not.

Eight

Associations

In CML, an *association* represents a relation between two *concepts* (§4), where a reference to an *instance* of each *concept* is found in every tuple that is part of the relation. When *concepts* have an *association* between themselves, its *instances* are linked in such way that it is possible to access an *instance* of one *concept* from an *instance* of the other *concept*.

The UML [4] metamodel has a metaclass named *Association* that has *Property* instances, whose *types* are the *Class* instances that are part of the *association*. In UML, the name of each *Property* instance in the *Association* metaclass is known as the *role* of the corresponding *Class* in the *association*.

On the CML metamodel, on other hand, the *Association* metaclass is only needed when it is necessary to define *bidirectional associations*, whose *links* are accessible from either *association end*. For *unidirectional associations*, where only one *association end* is accessible, only a *property* is defined in the source *concept*, making its *type* the target *concept*.

On the ER [1] metamodel, each *association* is known as a *relationship set*, and each tuple in this set is called a *relationship*. Unlike CML and UML, the tuples in a *relationship set* of an ER model can be queried directly, and no notion of *property* is required as part of the *entity type* in order to access those *relationships*.

As it is case for *attributes* (§6), *associations* in CML can also be derived from other *associations* (just as well as in UML); they are called *derived associations* (§9).

8.1 Example

Listing 8.1 presents some examples of *associations* declared in CML. The concept **Vehicle** contains the property **driver**, which may optionally refer to an instance of **Employee**, meaning that a **driver** may or may not be assigned to a single **Vehicle**. The concept **Vehicle** also has the property **owner**, which always refers to an instance of **Organization**, meaning that an **owner** must always be assigned to each instance of **Vehicle**. Similarly, the concept **Employee** has the property **employer**, which must always be assigned to an instance of **Organization**.

Just below the declaration of **Organization**, we observe an association named **Employment**, which enumerates two *properties*: the first is **employer** from the concept **Employee**; the second is **employees** from the concept **Organization**. What this *association* implies is a correspondence between these two properties. Every time a reference to an instance of **Organization** is assigned to the slot **employer** of an instance of **Employee**, a reference to this same instance of **Employee** must be assigned to the slot **employees** of the **Organization** instance. However, since the *type* of **employees** in the concept **Organization** is a sequence (§21) of **Employee** instances, the reference to the instance of **Employee** will actually be appended to the sequence being held by the slot **employees** of the concept **Organization**, and maintained along with the other **Employee** instances already found in the sequence. Thus, the association **Employment** actually characterizes a *bidirectional association*.

The association **VehicleOwnership** is another example of a *bidirectional association*; in this case, between **Vehicle**'s **owner** property and **Organization**'s **fleet** property. It can be noticed, though, in this second *bidirectional association*, that the *types* of the *properties* are declared along with their names; such a *type* declaration, in the *association* declaration, is optional in CML, but must match the original *property* declaration under the *concept* declaration, if present.

The **driver** property in the concept **Vehicle** is a different case, since this *property* does not participate in any *association* declaration on listing 8.1. That's because there is no corresponding *property* in the concept **Employee** representing the other end of the *association*. As such, the property **driver** is representing the source end of a *unidirectional association*.

The property **drivers** in the concept **Organization** is *derived association* (§9).

```
@concept Vehicle
{
    plate: String;
    driver: Employee?;
    owner: Organization;
}

@concept Employee
{
    name: String;
    employer: Organization;
}

@concept Organization
{
    name: String;
    employees: Employee*;
    fleet: Vehicle*;
    /drivers = fleet.driver;
}

@association Employment
{
    Employee.employer;
    Organization.employees;
}

@association VehicleOwnership
{
    Vehicle.owner: Organization;
    Organization.fleet: Vehicle*;
}
```

Listing 8.1: Association Example

```

associationDeclaration
  returns [ Association association ]:
    'association ' NAME
    '{ ' (associationEndDeclaration ' ; ')* ' } ';

associationEndDeclaration
  returns [ AssociationEnd associationEnd ]:
    conceptName=NAME ' . ' propertyName=NAME
    ( ': ' typeDeclaration )?;

```

Listing 8.2: Association Concrete Syntax

8.2 Syntax

The concrete syntax used to declare an *association* in CML is specified by listing 8.2. First, the **association** keyword is followed by a NAME. Then, a list of *association ends* are declared under the **association** block. For each declaration of an *association end*, The **conceptName** and **propertyName** are optionally followed by a **typeDeclaration**.

The *Association* metaclass is presented in the EMOF [5] class diagram of figure ??, and its instantiation from the concrete syntax is specified by listing 8.3. For each parsed *association*, an instance of the *Association* metaclass will be created, and its meta-properties will be assigned according to parsed information:

- *name*: assigned with the value of the token NAME.
- *members*: an *ordered set* referencing all *associationEnd* instances parsed in the **association** block.

8.3 Constraints

The invariants of the metaclasses *Association* and *AssociationEnd* are specified by listing 8.4:

- *association_end_property_found_in_model*: Each *association end* enumerated under an *association* must correspond to a *property* of the same name under the specified *concept*.

```

node Association :
  'association' NAME '{' ( AssociationEnd ';' ) * '}'
{
  name = NAME;
  members = AssociationEnd *;
}

node AssociationEnd :
  conceptName=NAME '.' propertyName=NAME ( ':' type=Type )?
{
  concept = Model.concepts
    ->select( concept.name = conceptName )
    ->first()

  property = concept.allProperties
    ->select( property.name = propertyName )
    ->first()
}

```

Listing 8.3: Association AST Instantiation

- *association_end_type_matches_property_type*: If a *type* is specified for a given *association end*, its *name* and *cardinality* must match the *type* of the corresponding *property*.
- *association_must_have_two_association_ends*: An *association* must have exactly two *association ends*, since only *binary associations* are supported in CML.
- *association_end_types_must_match*: The *concept* of one *association end* must correspond to the *type* of the *property* of the other *association end*, and vice-versa.

8.4 Translation

```

context AssociationEnd
inv association_end_property_found_in_model:
    concept->notEmpty() and property->notEmpty()

context AssociationEnd
inv association_end_type_matches_property_type:
    self.propertyType->notEmpty() and
    self.property->notEmpty() implies
        self.propertyType.name = property.type.name and
        self.propertyType.cardinality = property.type.cardinality

context Association
inv association_must_have_two_association_ends:
    associationEnds->count() = 2

context Association
def: first = associationEnds->first()
def: last = associationEnds->last()
inv association_end_types_must_match:
    associationEnds->count() = 2 and
    first->notEmpty() and last->notEmpty() and
    first.concept->notEmpty() and
    first.property->notEmpty() and
    last.concept->notEmpty() and
    last.property->notEmpty() implies
        first.concept.name = last.property.type.name and
        last.concept.name = first.property.type.name

inv property_must_be_part_of_single_association:

inv no_associations_of_primitive_types:

```

Listing 8.4: Association Constraints

Nine

Derived Associations

Ten

Generalization / Specialization

A *concept* (§4) in CML may be generalized by another *concept*. In other words, a *concept* may be considered a specialization of another *concept*. Generalized *concepts* have *properties* (§5) that apply to a larger set of instances, while specialized *concepts* have *properties* that only apply to a subset of those instances.

In the UML [4] metamodel, such generalization/specialization relationship between *classes* is known as *generalization*, which is the name of the metaclass in the UML metamodel. The original version of the ER [1] metamodel lacked this kind of relationship between *entity types*.

10.1 Example

Figure 10.1 presents some examples of generalization/specialization relationships declared in CML. As shown, a *concept* (§4) may specialize zero or more other *concepts*. The latter are called the generalizations, while the former is called the specialization. A generalization, such as **Shape**, may define *attributes* (§6), such as **color** and **area**, or also the *roles* in *unidirectional associations* (§8). Both *attributes* and *roles* are *properties* (§5) shared among all its specializations. Some of these *properties* may be redefined by the some of the specializations, as it is the case with the *area* property, which is redefined by **Rectangle**, **Rhombus** and **Square**. Some specializations may also define new *properties*, such as **width** and **height** in **Rectangle**, which characterize only instances of this specialization. A *concept* may be a specialization of two or more other *concepts*, as seen with **Square**, which specializes both **Rectangle** and **Rhombus**, and thus can redefine *properties* of both generalizations. If a *property* has been defined by more than one generalization, then it must be redefined by the

specialization in order to resolve the definition conflict, which is the case with **area** in **Square**. If a redefinition suitable for both generalizations is unattainable, it may be an indication that either the specialization or the generalizations are unsound from the domain's prospective.

10.2 Syntax

Figure ?? specifies the syntax used to declare a *concept* (§4), and in turn its generalizations. A list of NAMES may be enumerated after the declared *concept's* NAME, referring to other *concepts* that this concept is a specialization of.

Figure 4.1 presents the *Concept* metaclass in an EMOF [5] class diagram of the CML metamodel, and figure ?? specifies the *concept* transformation from its concrete syntax to its abstract syntax. There is a unidirectional association in the *Concept* class that keeps track of the generalization/specialization relationships, which is named *generalizations*. It is an *ordered set* referencing all *concepts* whose NAMES were enumerated in the *GeneralizationList* of the declared *concept*.

10.3 Constraints

Figure 10.2 presents the invariants of the *Concept* and *Property* classes related to *generalizations*:

- *not_own_generalization*: A *concept* (§4) may not be listed on its own *GeneralizationList*, nor on the *GeneralizationList* of its direct or indirect generalizations.
- *compatible_generalizations*: The *generalizations* of a *concept* must all be compatible between themselves, that is, no two *generalizations* may have a *property* with the same name but a different type.
- *generalization_compatible_redefinition*: A *property* may only be redefined with the same type defined in the *generalizations*.
- *conflict_redefinition*: A *concept* is required to redefine a *property* that has been defined by two or more of its *generalizations* in order to resolve the definition conflict. That is required only if the *property* has been initialized or derived in at least one of the *generalizations*. Otherwise, the redefinition is not required.

```

— Generalization of Circle and Rectangle:
@concept Shape
{
    — Specializations below share the color attribute as-is:
    color: String;

    — Specializations below redefine the area attribute:
    area: Double;
}

— Specialization of Shape:
@concept Rectangle: Shape
{
    — New attributes that characterize a rectangle:
    width: Double;
    height: Double;

    — Redefinition of the area attribute:
    /area = width * height;
}

— Another specialization of Shape:
@concept Rhombus: Shape
{
    — Diagonal attributes that characterize a rhombus:
    p: Double;
    q: Double;

    — Another redefinition of the area attribute:
    /area = (p * q) / 2.0d;
}

— Specialization of both Rectangle and Rhombus:
@concept Square: Rectangle, Rhombus
{
    — Only attribute needed to characterize a square:
    side_length: Double;

    — Redefinitions of Rectangle's attributes:
    /width = side_length;
    /height = side_length;

    — Redefinitions of Rhombus' attributes:
    /p = side_length * 1.41421356237d; — square root of 2
    /q = p;

    — Required to redefine area in order to resolve conflict
    — between Rectangle's area and Rhombus' area:
    /area = side_length ^ 2.0d;
}

```

Figure 10.1: Generalization Examples

```

context Concept:: all_generalizations: Set(Concept)
derive:
    generalizations -> closure (generalizations)

context Concept:: all_properties: Set(Property)
pre:
    all_generalizations -> excludes (self)
derive:
    elements -> union (
        generalizations.all_properties -> select (p1 |
            not elements -> exists (p2 | p1.name == p2.name)
        )
    )

context Concept:: generalization_pairs
    : Set(Tuple(left: Concept, right: Concept))
derive:
    generalizations -> collect (g1 |
        generalizations
        -> select (g2 | g1 != g2)
        -> collect (g2 | Tuple { left: g1, right: g2 })
    ) -> flatten ()

context Concept:: generalization_property_pairs
    : Set(Tuple(left: Property, right: Property))
derive:
    generalization_pairs -> collect (pair |
        pair.left.all_properties -> collect (p1 |
            pair.right.all_properties
            -> select (p2 | p1 != p2 and p1.name = p2.name)
            -> collect (p2 | Tuple { left: p1, right: p2 })
        ) -> flatten ()
    ) -> flatten ()

context Concept
inv not_own_generalization:
    all_generalizations -> excludes (self)

context Concept
inv compatible_generalizations:
    generalization_property_pairs
    -> forAll (
        left.type.name = right.type.name and
        left.type.cardinality = right.type.cardinality
    )

context Concept
inv conflict_redefinition:
    generalization_property_pairs
    -> select (left.type = right.type)
    -> select (left.derived or left.expression -> notEmpty () or
        right.derived or right.expression -> notEmpty ())
    -> forAll (self.elements -> exists (name = left.name))

context Property
inv generalization_compatible_redefinition:
    self.scope.generalizations.all_properties
    -> select (property | self.name = property.name)
    -> forAll (property |
        self.type.name = property.type.name and

```

Eleven

Abstract Concepts

An *abstract concept* is one that does not represent specific instances, but instead serves as a *generalization* (§ 10) for other *concepts*, which in turn represent specific instances. Thus, all instances of an *abstract concept* are first instances of its *specializations*. CML supports tagging a *concept* as *abstract*.

An *abstract concept* in CML may also define a *derived property* (§ 5) without providing an *expression* (§ 12) in its definition; such *properties* may also be called *abstract properties*.

CML's support for *abstract concepts* matches UML's [4], which allows the declaration of *abstract classes* – by setting the *isAbstract* attribute of the *Class* metaclass instance to *true*. UML also allows the declaration of corresponding *abstract attributes* and *abstract operations*.

The original version of the ER [1] metamodel, however, as a consequence of lacking the *generalization/specialization* relationship, has not considered the notion of *abstract entities*.

11.1 Example

Listing 11.1 presents an example of an *abstract concept* declared in CML. As shown, the concept **Shape** is tagged as *abstract*, and as such no direct instances of *Shape* are ever instantiated. As an *abstract concept*, **Shape** can define *abstract properties*, like **area**, which is just a *derived property* (§ 5) without an *expression* (§ 12). An *abstract concept* may also define concrete *properties*, such as **color** in **Shape**. The concept **Circle** is a *specialization* of **Shape** that must redefine the property **area** (and provide an *expression*) if it is to be considered a *concrete concept*. As a *concrete concept*, **Circle** may have direct instances, which are in turn instances of *Shape* as well. **Circle**

may also redefine *concrete properties* of **Shape**, like **color**, but the redefinition is not a requirement in this case. In **UnitCircle**, we can observe that the redefinition of an *abstract property*, such as **area**, may be made *concrete*; meaning it does not need to be redefined as a *derived property*. The converse situation is also allowed in CML, where a *concrete property* is redefined by as a *derived property*, as illustrated with the property **radius** in **UnitCircle**.

11.2 Syntax

Listing 4.2 specifies the syntax used to declare a *concept* (§4) in CML. It shows that a *concept* may be tagged with the **abstract** keyword in order to convey it as an *abstract concept*. Listing 5.2 specifies the syntax used to declare a *property* (??) in CML. It shows that a *property* may be prefixed with a forward slash ("/") in order to mark it as a *derived property*. If the optional **expression** is not provided, the property is then considered an *abstract property*.

Figure 4.1 presents the *concept* metamodel in an EMOF [5] class diagram, and listing 4.3 specifies the *concept* transformation from its concrete syntax to its abstract syntax. There is a **Boolean** attribute named **abstract** in the *Concept* class that determines whether a *concept* is *abstract* or not.

11.3 Constraints

Listing 11.2 presents the invariants of the *Concept* and *Property* classes in CML's EMOF [5] metamodel related to *abstract concepts*:

- *abstract_property_redefinition*: A *concrete concept* must redefine concretely all *abstract properties* of its *generalizations*.
- *abstract_property_in_abstract_concept*: Only *abstract concepts* may have *abstract properties*.

```

— As an abstract concept,
— no direct instances of Shape are ever created.
@abstraction Shape
{
  — A derived property without an expression
  — is considered abstract.
  — Only abstract concepts may have abstract properties.
  /area: Double;

  — Abstract concepts may also have concrete properties:
  color: String;
}

— All instances of Circle are in turn instances of Shape.
@concept Circle: Shape
{
  radius: Double;

  — In order to be considered a concrete concept,
  — Circle must redefine the abstract properties
  — inherited from Shape.
  /area = 3.14159d * radius ^ 2;

  — Circle may also redefine concrete properties of Shape.
  — However, the redefinition is not required in this case.
  color = "Blue";
}

@concept UnitCircle: Circle
{
  — Observe below that the redefinition of
  — an abstract property may be concrete;
  — that is, it does not have to be derived
  — as it was done in Circle.
  area = 3.14159d;

  — In the case above, however,
  — it is desirable to redefine "area" as a derived property,
  — in order to guarantee area's value cannot be modified
  — after the instantiation of UnitCircle.
  — This is done with the redefinition of "radius" below.
  — Notice that, in Circle, radius was concrete,
  — but its redefinition below makes it derived.
  — That's allowed in CML just as the other way around,
  — as it was done with "area" above.
  /radius = 1.0d;
}

```

Listing 11.1: Abstract Concept Example

```
context Property :: abstract: Boolean
derive:
  self.derived and self.expression -> isEmpty()

context Property :: concrete: Boolean
derive:
  not self.abstract

context Concept
inv abstract_property_redefinition:
  self.concrete implies
    self.generalizations.all_properties
      ->select(abstract)
      ->forall(p1|
        self.properties
          ->select(p2| p1.name = p2.name)
          ->reject(abstract)
          ->notEmpty()
      )

context Property
inv abstract_property_in_abstract_concept:
  self.abstract implies self.scope.abstract
```

Listing 11.2: Abstract Concept Constraints

Part III

Expressions

Twelve

Expressions

An *expression* in CML is used to compute *values* (§20) and sequences (§21) that initialize or derive *properties* (§5). On the UML [4] metamodel, it corresponds to the *Expression* metaclass; in OCL [3], to *OclExpressionCS*.

CML *expressions* are designed to provide the same level of expressivity provided by OCL *expressions*, but the CML syntax varies from OCL's; especially for OCL's *collection operations*, which correspond to *query expressions* (§18) in CML.

12.1 Example

Listing 12.1 has some examples of CML *expressions*. As shown, there are different types of expressions: *literal values* (§13), *prefix expressions* (§14), *infix expressions* (§15), *conditional expressions* (§16), *path expressions* (§17) and *query expressions* (§18).

12.2 Syntax

Listing 12.2 specifies the syntax of all kinds of *expressions* in CML. It also lists them in their order of precedence.

Observe that the grammar in listing 12.2 has left recursions (and thus it is ambiguous). However, ANTLR [6] resolves the ambiguity based on the order in which the *expression* alternatives are listed, and so the order in listing 12.2 defines the precedence among the operators for CML.

Also, according to ANTLR, and as required by CML, all *expressions* in the grammar are left-to-right associative, except for the *exponentiation expression*, which is right-to-left associative, as defined by the **<assoc=right>** clause.

```

@concept Expressions
{
  — Literal Values:
  c: String = "SomeString";
  d: Integer = 123;

  — Prefix Expression:
  minus_sign = -2;

  — Infix Expressions:
  addition = 1 + 2;
  equality = 3 == 3;
  boolean_expr = q and p;

  — Conditional Expression:
  if_then_else = if a > 0 then a else b;

  — Path Expression:
  path = somePath.bar;

  — Query Expression:
  select_query = items | select: name == "this";
}

```

Listing 12.1: Expression Example

Figure 12.1 presents the *Expression* metaclass in an EMOF [5] class diagram. For each kind of *expression* parsed by the compiler, an instance of an *Expression* subclass will be created, and its properties will be assigned according to parsed information:

- *kind*: a *String* value matching the *Expression* subclass; for example, for the *Literal* subclass, **kind = "literal"**.
- *type*: a derived attribute that computes the *Type* of the *expression*; each *Expression* subclass will do its own *Type* computation by providing its own definition for this derived attribute.

```

expression returns [Expression expr]
: literalExpression
| pathExpression
| operator=('+' | '-' | NOT) expression
| <assoc=right> expression operator='^' expression
| expression operator=('*' | '/' | '%') expression
| expression operator=('+' | '-') expression
| expression operator=('<' | '<=' | '>' | '>=') expression
| expression operator=('==' | '!=') expression
| expression operator=AND expression
| expression operator=OR expression
| expression operator=XOR expression
| expression operator=IMPLIES expression
| IF cond=expression
  THEN then=expression
  ELSE else_=expression
| queryExpression
| '(' inner=expression ')';

queryExpression returns [Expression expr]
: pathExpression
| joinExpression
| queryExpression '|' transformDeclaration;

joinExpression returns [Join join]:
  FOR enumeratorDeclaration (',' enumeratorDeclaration)*;

enumeratorDeclaration:
  var=NAME IN pathExpression;

transformDeclaration returns [Transform transform]:
  (FROM var=NAME '=' init=expression)?
  operation=
    ( SELECT      | REJECT
    | YIELD        | RECURSE
    | INCLUDES     | EXCLUDES
    | EVERY        | EXISTS
    | REDUCE
    | TAKE         | DROP
    | FIRST        | LAST
    | COUNT        | SUM          | AVERAGE
    | MAX          | MIN
    | REVERSE)
  suffix=(UNIQUE | WHILE)?
  expr=expression?;

```

Listing 12.2: Expression Concrete Syntax

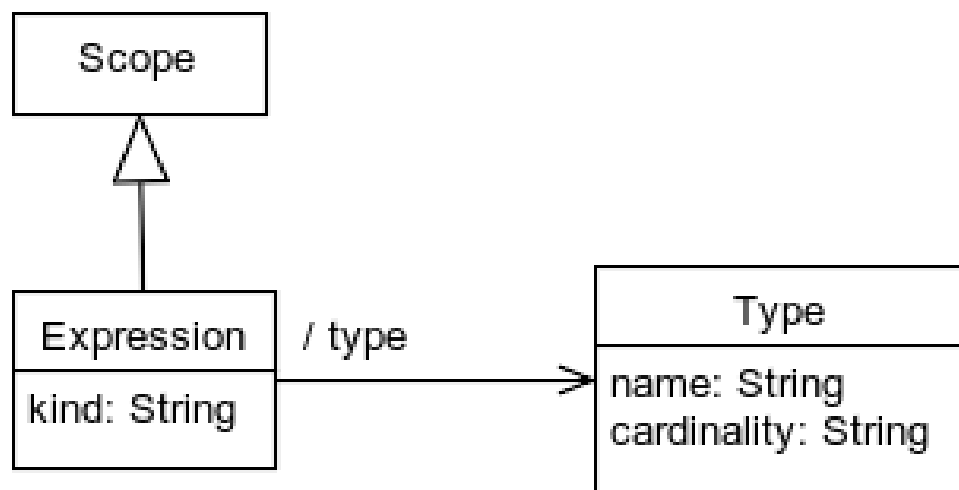


Figure 12.1: Expression Abstract Syntax

Thirteen

Literal Values

Fourteen

Prefix Expressions

Fifteen

Infix Expressions

Sixteen

Conditional Expressions

Seventeen

Path Expressions

Eighteen

Query Expressions

Part IV

Type Checking

Nineteen

Types

Twenty

Primitive Types

A *primitive type* in CML is one of the pre-defined *data types* supported by the language, as shown in tables 20.1 and 20.2.

In the ER [1] metamodel, a *data type* is formally defined as a *set of values* that can be held by an *attribute* (§6). The original ER paper [1] states that, for each *value set* (i.e. *data type*), there is a *predicate* that can be used to test whether a *value* belongs to the *set*. In CML, instead, *literal expressions* are syntactically defined for each *primitive type*, so that the *type* can be inferred from the *literal expression*.

On the original ER paper, it is also said that *values* in a *value set* may be equivalent to *values* in another *value set*. In CML, also, *literal expressions* of the *Integer* type may be equivalent to *literal expressions* of the *Decimal*, and so with other *numeric types*. This allows *expressions* of a *primitive type* to be promoted to *expressions* of another *primitive type* in order to allow *type inference* of composite *expressions*, such as *infix expressions* (§15).

In the UML [4] metamodel, there is a specific metaclass named *PrimitiveType*, which matches to the same notion in CML.

20.1 Example

Figure 20.1 presents examples of *attributes* declared with *primitive types* in CML. Each example corresponds to one of the *primitive types* supported by the language, as shown in tables 20.1 and 20.2. The *target constructors* (§23) of CML's base module will translate the *primitive types* to Java, C#, C/C++, Python, and TypeScript (JavaScript), according to the mapping shown in the tables.

```
@concept PrimitiveTypes
{
    — Core Primitive Types:

    — Only values are the literal expressions: true, false
    a: Boolean;

    — 32-bit signed two's complement integer
    c: Integer;

    — Arbitrary precision arithmetic.
    — BigDecimal in Java; decimal in C#; decimal128 in C++.
    d: Decimal;

    — 16-bit Unicode character sequences
    — as in Java, C#, C++ (std::wstring), and JavaScript.
    b: String;

    — Additional Primitive Types:

    — 8-bit signed two's complement integer
    e: Byte;

    — 16-bit signed two's complement integer
    f: Short;

    — 64-bit signed two's complement integer
    g: Long;

    — 32-bit IEEE 754 floating point
    h: Float;

    — 64-bit IEEE 754 floating point
    i: Double;
}
```

Figure 20.1: Example of *Primitive Types*

CML	Java	C#	C++	Python	TypeScript (JavaScript)
String	String	string	std::wstring	str	string
16-bit Unicode character sequences.					
Boolean	boolean	bool	bool	bool	boolean
Only values are the literal expressions: true , false .					
Integer	int	int	int32_t	int	number
32-bit signed two's complement integer.					
Decimal*	BigDecimal	decimal	decimal128	Decimal	number
Arbitrary precision, fixed-point, or decimal floating-point, depending on the target language.					

*The specification of Decimal type varies by target programming language. Compared to the binary floating-point types (Float and Double), the Decimal type is better suited for monetary calculations at a performance cost.

Table 20.1: Core Primitive Types in CML.

```
typeDeclaration returns [Type type]:  
    NAME cardinality ?;  
  
cardinality:  
    ('?' | '*');
```

Figure 20.2: Type Declaration Syntax

20.2 Syntax

Figure 20.2 specifies the syntax used to declare any kind of *type*, including *primitive types*. The NAME of the *type* may be any of the *primitive types* defined in the column named CML of the tables 20.1 and 20.2. Optionally, cardinality may also be specified for a *primitive type*. The “*” cardinality suffix allows zero or more values to be stored in a property as a collection type (??). The “?” cardinality suffix allows a single value to be stored, or none. If no cardinality is specified, a value must be assigned to the *attribute* when its *concept* is instantiated.

CML	Java	C#	C++	Python	TypeScript (JavaScript)	Specification
Byte	byte	byte	int8_t	int	number	8-bit signed two's complement integer
Short	short	short	int16_t	int	number	16-bit signed two's complement integer
Long	long	long	int64_t	long	number	64-bit signed two's complement integer
Float	float	float	float*	float	number	32-bit IEEE 754 binary floating point
Double	double	double	double*	float	number	64-bit IEEE 754 binary floating point

*C++ floating point types may vary by hardware and compiler

Table 20.2: Additional Primitive Types in CML.

```
node Type: NAME CARDINALITY?  
{  
  name = NAME;  
  cardinality = CARDINALITY?;  
}
```

Figure 20.3: Type AST Instantiation

Figure ?? presents the *Type* metaclass in an EMOF [5] class diagram of the CML metamodel, and figure 20.3 specifies the transformation from the *type* concrete syntax to its abstract syntax.

Twenty-one

Sequence Types

Part V

Code Generation

Twenty-two

Templates

Twenty-three

Constructors

Twenty-four

Tasks

Twenty-five

Targets

Part VI

Organization and Sharing

Twenty-six

Modules

Twenty-seven

Libraries

Part VII

Appendices

A

CML Concrete Syntax (Grammar)

A.1 ANTLR Grammar

```

// Compilation Units:

compilationUnit:
    declarations*;

declarations:
    moduleDeclaration | conceptDeclaration | associationDeclaration | taskDeclaration;

// Concept Declarations:

conceptDeclaration returns [Concept concept]:
    ABSTRACT? 'concept' NAME
    (':' generalizations)?
    ( ';' | propertyList);

generalizations:
    NAME (',' NAME)*;

// Property Declarations:

propertyList:
    '{' (propertyDeclaration ';')* '}';

propertyDeclaration returns [Property property]:
    DERIVED? NAME (':' typeDeclaration)? ('=' expression)?;

DERIVED: '/';

// Type Declarations:

typeDeclaration returns [Type type]:
    NAME cardinality?;

cardinality:
    ('?' | '*');

```

```
// Target Declarations:

targetDeclaration returns [Target target]:
    'target' NAME propertyList;

// Names:

// All keywords must be declared before NAME.
// Otherwise, they are recognized as a NAME instead.

FOR: 'for';
IN: 'in';

SELECT: 'select';
REJECT: 'reject';

YIELD: 'yield';
RECURSE: 'recurse';

INCLUDES: 'includes';
EXCLUDES: 'excludes';

EVERY: 'every';
EXISTS: 'exists';

FROM: 'from';
REDUCE: 'reduce';

TAKE: 'take';
DROP: 'drop';

FIRST: 'first';
LAST: 'last';

COUNT: 'count';
SUM: 'sum';
AVERAGE: 'average';
MAX: 'max';
MIN: 'min';
REVERSE: 'reverse';
```

```

UNIQUE: 'unique';
WHILE: 'while';

IF: 'if';

THEN: 'then';

ELSE: 'else';

BOOLEAN: 'true' | 'false';

AND: 'and';

OR: 'or';

XOR: 'xor';

IMPLIES: 'implies';

NOT: 'not';

ABSTRACT:
    'abstract';

NAME:
    ('A'..'Z' | 'a'..'z')
    ( 'A'..'Z' | 'a'..'z' | '0'..'9' | '_' )*;

// Literals:

literalExpression returns [Literal literal]: BOOLEAN | STRING | INTEGER | LONG | SHORT

STRING:
    '"' (ESC | . )*? '"';

fragment ESC: '\\'[btnr"\\];

INTEGER:
    ('0'..'9')+;

```

```
LONG:
    ('0'..'9')+ 'l';

SHORT:
    ('0'..'9')+ 's';

BYTE:
    ('0'..'9')+ 'b';

DECIMAL:
    ('0'..'9')* '.' ('0'..'9')+;

FLOAT:
    ('0'..'9')* '.' ('0'..'9')+ 'f';

DOUBLE:
    ('0'..'9')* '.' ('0'..'9')+ 'd';

// Ignoring Whitespace:

WS:
    ( ' ' | '\t' | '\f' | '\n' | '\r' )+ -> skip;

// Ignoring Comments:

COMMENT:
    (('//' | '-') .*? '\n' | '/*' .*? '*/' ) -> skip;
```


B

CML Abstract Syntax (Metamodel)

C

CML Abstract Syntax Tree (Instantiation)

D

CML Constraints (Validations)

```
context Concept

inv unique_concept_name:
  parent.concepts
    ->select(c| c != self and c.name = self.name)
    ->isEmpty()

context Property
inv unique_property_name:
  self.scope.properties
    ->select(p| p != self and p.name = self.name)
    ->isEmpty()

context Property
inv property_type_specified_or_inferred:
  type->notEmpty() or expression->notEmpty()

context Property
inv property_type_assignable_from_expression_type:
  type->notEmpty() and expression->notEmpty() implies
    type.isAssignableFrom(expression.type)
```

E

Language Specification Notation

Bibliography

- [1] Peter Pin-Shan Chen. The Entity-Relationship Model (Reprinted Historic Data). In David W. Embley and Bernhard Thalheim, editors, *Handbook of Conceptual Modeling: Theory, Practice, and Research Challenges*, pages 57–84. Springer Berlin Heidelberg, Berlin, Heidelberg, 2011.
- [2] Torben Ærgidius Mogensen. *Introduction to Compiler Design*. Undergraduate Topics in Computer Science. Springer, 2011.
- [3] OMG. Object Constraint Language (OCL), Version 2.4, 2014.
- [4] OMG. Unified Modeling Language (UML), Superstructure, Version 2.5, 2015.
- [5] OMG. Meta Object Facility (MOF) Core Specification, Version 2.5.1, 2016.
- [6] Terence Parr. *The Definitive ANTLR 4 Reference*. Pragmatic Bookshelf, 2nd edition, 2013.
- [7] Terence John Parr. Enforcing Strict Model-view Separation in Template Engines. In *Proceedings of the 13th International Conference on World Wide Web, WWW '04*, pages 224–233. ACM, New York, NY, USA, 2004.